

SEP Engine: A Quantum-Inspired Manifold Architecture for Emergent Pattern Analysis

Abstract

The **Self-Emergent Processor (SEP) Engine** is a high-performance, quantum-inspired framework for real-time analysis of complex data streams (e.g. financial markets). It quantifies **informational coherence** and **stability** within a dynamic data manifold to discern predictive patterns and emergent forces in chaotic environments ¹. The engine's architecture separates a **reality-tracking core** – a GPU-accelerated processing manifold backed by a Valkey database – from a lightweight **decision/rendering interface**. The Valkey key-value store serves as a **time-anchored manifold of identities**, where each identity's state at a given timestamp is stored with its key metrics (price, entropy, stability, coherence). These metrics form a **3D surface** in the information space that evolves and **settles over time** as patterns emerge and stabilize. A remote GPU kernel continuously updates this “hot band” of recent keys via deterministic recomputation of metrics, ensuring that state transitions are fully reproducible ². Owing to deterministic mappings between successive states, prior identity states can be derived or verified from present observables. The front-end UI reads from the manifold in real-time (via an API) to visualize the **current stability** and coherence of the system, but performs no computations itself. This modular architecture cleanly **divides reality-tracking from decision-rendering**, allowing theoretical analysis of the manifold dynamics (for researchers in quantum computing, statistical inference, and non-Euclidean optimization) while providing a robust operational interface for practitioners.

1. Introduction

Modern complex systems – from financial markets to neural data – exhibit emergent patterns that challenge classical modeling. The **SEP Engine** addresses this by treating data streams as evolving manifolds of information. It was designed as a modular C++ engine with CUDA acceleration, implementing algorithms inspired by quantum information theory ³. In essence, SEP quantifies how *coherent* and *stable* a pattern is within noisy data, enabling detection of self-organizing structure beyond traditional statistical signals ¹. By leveraging **quantum-inspired algorithms** – notably *Quantum Fourier Hierarchy* (QFH) and *Quantum Bit State Analysis* (QBSA) – the engine probes phase alignment and coherence in the data ³. These algorithms simulate “emergent informational forces” within the data manifold to identify stable patterns ³. The result is a system that operates at **quantum-like speed on classical hardware** ⁴, offering real-time insight into pattern evolution and predictability, with direct applications in quantitative trading and beyond.

Crucially, the SEP Engine's design emphasizes a clear **separation of concerns**. The heavy-lifting of state evolution and pattern computation is handled by a GPU-accelerated core, while the user interface and external clients interact through a high-level API ⁵. This paper presents the current architecture of SEP, focusing on how the **Valkey database** underpins a time-indexed identity manifold, how **coherence**, **stability**, and **entropy** metrics define a dynamic information surface, and how the **deterministic GPU kernel** drives state transitions. We also highlight the theoretical implications of this design, such as determinism enabling backward inference and the non-Euclidean geometry of the metric space, making

SEP of interest to researchers in quantum computing and statistical inference. Throughout, references to the codebase (documentation and source) are provided to ground each aspect in the implemented system.

2. System Architecture Overview

Core Structure: The SEP Engine is delivered as a single compiled binary (`sep_engine`) composed of cohesive C++ static libraries ⁶. Key modules include `libsep_core` (foundational services and logging), `libsep_memory` (tiered memory and data caching), `libsep_quantum` (analytical kernel algorithms), and `libsep_api` (external interface server) ⁷ ⁸. The design enforces unidirectional dependencies: for example, the high-level API orchestrates the quantum and memory modules, which in turn rely on core utilities ⁹ ⁵. This ensures a modular architecture where each subsystem has a clear role. The **GPU acceleration** is abstracted via `libsep_compat`, which manages CUDA device resources and kernel launches ¹⁰. In effect, the GPU acts as a “supercharger” for the engine’s computations, massively increasing throughput for pattern analysis ¹¹. The engine’s primary output is a composite **predictive gauge** metric that combines coherence, stability, and entropy measures ¹², though internally these components are tracked separately (see §3).

Reality-Tracking Manifold vs UI: A fundamental split in the architecture is between the **reality-tracking engine** and the **decision/rendering layer**. The engine’s core (including the GPU kernels and memory hierarchy) is responsible for ingesting data, evolving patterns, and updating the state manifold in real-time. The **Valkey database** – a high-performance in-memory key-value store – is deployed as part of this core, functioning as the engine’s long-term memory (LTM) backend ¹³. The front-end (UI or any external client) does not perform calculations; it interfaces through `libsep_api` to retrieve the latest metrics and state information ⁵. This means the **UI simply visualizes** the system’s current state (e.g. plotting stability/coherence gauges) without computing those values itself. By isolating computation in the back end and visualization in the front end, SEP achieves both high performance and clarity: the engine can track reality at full throttle, while the operator’s interface remains responsive and focused on interpretation. Notably, even specialized visualization modules (e.g. a Blender 3D rendering plugin) are kept separate from the core logic, interfacing only through defined APIs ¹⁴. This clean separation aligns with best practices in system design and is pivotal for scalability and testability of the engine ¹⁵.

3. Time-Anchored Identity Manifold (Valkey Database)

At the heart of SEP’s architecture is the **Valkey database**, which serves as a **time-anchored manifold of identities** in the system. In practical terms, Valkey (a high-performance key-value store akin to an advanced Redis) is used as the **Long-Term Memory (LTM)** tier to persistently store patterns that have proven stable and important ¹³. Each **identity** in this context refers to a distinct emergent pattern or data entity that the engine is tracking. Identities could represent, for example, a candidate market regime pattern or a learned signal feature. As time progresses, each identity’s state is captured at discrete intervals (e.g. each tick or analysis window) and stored as an entry in Valkey. Thus, the database becomes a chronologically indexed manifold: traversing keys in order is akin to walking along the temporal trajectory of each identity.

Key-Snapshot Mapping: Each **key** in Valkey corresponds to a fixed-time snapshot of an identity’s state. The value stored at that key encapsulates both raw data and computed metrics at that time. In the case of a trading application, for instance, a snapshot may correspond to one time-step of market data (e.g. one

OHLC bar or a window of ticks) and the pattern the engine has extracted from it. Concretely, each stored snapshot includes:

- **Timestamped Market Data** – e.g. the price or OHLC features for that time slice, serving as the raw input context for the pattern.
- **Entropy**: a measure of unpredictability or information content in the recent data, typically the normalized Shannon entropy of the pattern's distribution ¹⁶. A high entropy indicates a more disordered or noisy state, whereas lower entropy suggests more information structure (signal) in the pattern.
- **Coherence**: a dimensionless metric of the pattern's internal alignment and self-similarity. In implementation, coherence is calculated as an inverse variation (e.g. $\text{coherence} = 1 / (1 + CV)$ where CV is a coefficient of variation) ¹⁶. This means a perfectly regular pattern (low variance in its features) approaches coherence ≈ 1 , while a highly erratic pattern yields coherence closer to 0. Coherence effectively quantifies how well the identity “holds together” as a consistent shape or signal.
- **Stability**: a metric gauging the persistence of the pattern through time. This is computed in a variance-weighted manner ¹⁶, reflecting how much the identity's characteristics change from one state to the next. A stable identity shows minimal change (variance) over successive snapshots, indicating an enduring pattern, whereas an unstable one fluctuates rapidly.

Each identity's trajectory can thus be visualized as a **path through a 3D metric space** defined by (Coherence, Stability, Entropy) coordinates over time. In this space, time acts as an implicit fourth dimension anchoring the manifold. The collection of all identity states at a given time forms a surface in the metric space – essentially a cross-section of the manifold at that time slice. As the system runs, this surface is **dynamic**: new points (snapshots) are added and existing ones move or fade out as identities evolve or cease. Crucially, because identities that prove *persistent and coherent* are retained in the LTM (Valkey) store, the surface of points tends to **settle over time**. Early on, many tentative patterns (identities) may populate the surface with widely varying entropy and coherence. Many of these “noisy” identities dissipate or are not promoted to long-term storage if they fail to maintain stability. What remains (and gets reinforced in LTM) are the higher-coherence, higher-stability identities. In other words, the manifold naturally filters itself: the chaotic, high-entropy parts of the surface evaporate, leaving behind ridges or clusters of stable, coherent patterns. This self-organization can be seen as the surface **converging** or smoothing out. Mathematically, one could imagine an energy landscape over the manifold where coherence increases and entropy decreases as a pattern stabilizes – identities gravitate towards “basins” of high stability/coherence. By anchoring every state in time via the Valkey keys, the engine ensures that *every point on this informational surface is catalogued*. This provides an audit trail of the system's reality: any prior state of any identity can be retrieved by key lookup. In effect, **the Valkey manifold is a ledger of emergent pattern evolution**, one that is queryable in hindsight for research or verification.

4. Dynamic Metric Surface and Non-Euclidean Geometry

The triad of metrics (**entropy, stability, coherence**) defines a feature space that is inherently **non-Euclidean** in character. Each metric is derived from statistical or information-theoretic properties of the data, and they are not linearly independent in the way simple Euclidean coordinates are. For instance, coherence is a nonlinear function of variance (via $1/(1+CV)$), and entropy is a logarithmic measure of distribution spread. The interplay between these metrics creates a curved **information geometry**: small changes in an identity's data can have complex, interdependent effects on entropy and coherence.

Therefore, optimizing pattern quality (e.g. searching for a state of high coherence and stability but low entropy) is effectively an optimization problem on a curved manifold rather than in a flat vector space. This aligns with concepts in **information geometry** and **non-Euclidean optimization**, where one deals with metrics defined by statistical distances or divergences.

To illustrate, consider two pattern identities with equal coherence but different entropy – one might correspond to a clean periodic signal and another to a chaotic signal that by coincidence has similar self-similarity in a short window. Even though their coherence values are equal, their predictability differs, which the entropy captures. The *distance* (in terms of meaningful change) between these two states cannot be understood by coherence alone; one must consider the joint metric space. SEP's use of a **composite gauge** further underscores this geometric perspective: the engine often combines normalized coherence, stability, and entropy into a single **predictive gauge** value by a weighted sum ¹². These weightings effectively define a metric on the space of patterns (a particular linear slice through the non-Euclidean manifold) to rank which identities are most promising for prediction. The very form of the gauge, $\text{gauge} = w_c \cdot \text{coherence_norm} + w_s \cdot \text{stability_norm} - w_e \cdot \text{entropy_norm}$, treats increases in coherence or stability as positive contributions and entropy as a detractor ¹⁷. Tuning the weights (w_c, w_s, w_e) is akin to choosing a projection or metric on the manifold that best separates signal from noise for the domain at hand.

For theoretical researchers, this metric manifold offers a rich ground: one can apply concepts from **Riemannian geometry** or **information manifolds** to analyze how identities cluster or separate. Notably, SEP's metrics relate to well-studied quantities – e.g. Shannon entropy places identities on an information-theoretic landscape, while coherence relates to correlation structure (in quantum terms, akin to a measure of *pure state* vs *mixed state*). The **coherence register** concept used in SEP formalizes this: it implies that verification of a pattern's signal corresponds to projecting the pattern onto a latent basis and measuring alignment (coherence) ¹⁸. As patterns evolve, their successive projections form a Markov chain in this metric space, with **deterministic transitions** (since the projection operation is fixed) resulting in a smooth trajectory ¹⁹. Long-lived, coherent patterns correspond to long runs in the Markov chain that don't "break" – they represent a kind of stable attractor in the manifold. In contrast, unstable patterns quickly diverge (high entropy kicks in, coherence drops). This view connects with non-Euclidean optimization: finding an optimal pattern can be seen as searching for a path that climbs a "coherence hill" on a curved surface, a potentially NP-hard problem mitigated by SEP's specialized algorithms (like using sliding time windows and GPU parallelism for search) ²⁰. In summary, the SEP Engine provides not just raw metrics, but a geometrical lens – a dynamic surface where pattern discovery and verification are translated into movements and positions on a manifold, invoking deep connections to information geometry and even quantum state space analogies.

5. GPU-Accelerated Deterministic Computation

GPU Kernel Updates: The engine's core analytic work is performed by a dedicated GPU kernel service. SEP utilizes custom CUDA kernels (in `libsep_quantum`, e.g. `quantum_kernels.cu` and `pattern_kernels.cu`) to process incoming data and update pattern metrics in bulk ²¹ ⁸. This design exploits the parallelism of modern GPUs to handle the high-throughput requirements of real-time data streams. For example, as each new market tick or bar arrives, the engine can launch a GPU kernel that evaluates all active patterns against the new data, updating their coherence, stability, and entropy metrics simultaneously. This parallel update is crucial for maintaining a "hot band" of recent states in the manifold – essentially the leading edge of the timeline for each identity – with minimal latency. The **remote GPU** (which

can be on a server or cloud instance) acts as a continuously running co-processor for the engine, ingesting raw data and spitting out metric updates in a deterministic fashion. The rest of the system (memory tiers, database, etc.) is architected to accommodate these rapid updates, ensuring data moves efficiently from the GPU's output into the Valkey store or the appropriate memory tier.

Deterministic Recomputation: A key property of the GPU processing in SEP is that it is *deterministic*. Given the same input data and initial state, the GPU kernels will compute the same metrics and pattern evolution every time. There are no stochastic elements like random initializations or probabilistic branching in the core algorithms – even though the system is “quantum-inspired,” it is not quantum-random. This determinism has been intentionally designed and verified; one of the engine’s proof-of-concept milestones was **stateful and reproducible time-series analysis** ²². In practice, reproducibility means that if we replay historical market data through the engine (or if we checkpoint and restore the engine state), the sequence of identity states and their metrics will be identical ². This is extremely valuable for both debugging and theoretical analysis: it allows the **backward derivation of prior states** and cause-effect reasoning. Since every update function is known and invertible (at least in principle), one can take an observed state at time t and trace back to time $t-1$ by applying the inverse of the update function to the metrics, or simply by looking up the stored state at $t-1$ (since the manifold preserves it). In other words, the system’s evolution forms a deterministic chain through the state space, which can be walked in reverse. If an anomaly in coherence is observed now, the exact prior moment it arose can be pinpointed by checking when the metrics last changed, and because the metrics transition deterministically, we know that change was *caused* and not random. This property transforms the typically one-way street of time-series analysis into a two-way street: analysts can **project backward** to diagnose why a certain pattern became unstable or lost coherence by examining the precise metric deltas, with confidence that no hidden randomness is confounding the analysis.

From an implementation standpoint, the engine’s deterministic nature also comes from strict control of floating-point behavior and state management. The GPU kernels use fixed seeds or deterministic algorithms for any procedure that could otherwise introduce nondeterminism (e.g. they avoid non-associative reductions or use atomic operations carefully to ensure consistent results). Additionally, the promotion of patterns through memory tiers (STM → MTM → LTM) is done under fixed rules based on metric thresholds ²³. For instance, if coherence exceeds a configured threshold, the orchestrator will deterministically promote a pattern from short-term to medium-term storage ²⁴. These rules are set in configuration and yield the same outcome given the same metric inputs. The **MemoryTierManager** executes the data movement (allocate new block, copy, free old) without altering the data itself ²⁵, preserving the computed metrics exactly. Thus, the entire pipeline from data input, through GPU computation, to memory storage is deterministic and verifiable. This not only aids theoretical reproducibility but also ensures that the **Valkey manifold is consistent** – it never contains ambiguous or probabilistic entries, only definite states resulting from known transformations.

Backward Derivation of Identities: Thanks to this determinism and the time-anchored storage, one can achieve a form of *retrodiction* – inferring past states from present knowledge. In practice, if we know an identity’s state at time t , we typically also have stored states at times $t-1$, $t-2$, ... down to when the identity first appeared (because the engine promotes and retains stable identities). If for some reason an intermediate state was not stored (e.g. an identity was only promoted to LTM after it stabilized), the deterministic algorithm allows recomputing what the metrics *would have been* at earlier times by re-running the analysis on historical data. The combination of stored checkpoints and recomputation means the entire trajectory is reconstructable. This is reminiscent of reversible computing or time-reversible physical

simulations: given the end state and the rules of evolution, the path can be retraced. In SEP's case, the **mapping between successive identity states** is effectively one-to-one under fixed conditions, so *observing the metrics at time t provides strong constraints on what the metrics were at $t-1$* . For example, a big drop in coherence that is observed can be conclusively tied to the moment a particular entropy spike occurred in the data – by looking one step back, we see entropy was lower, and coherence was higher, etc., pinpointing the transition ¹⁹. This property elevates the SEP Engine from a black-box predictor to a **transparent dynamics simulator** of pattern evolution. It empowers researchers to validate theories (like how quickly should a coherent pattern decay in high entropy noise) directly against the engine's recorded state sequence, and to do so knowing the engine's results are perfectly repeatable and not flukes of random chance.

6. UI and Operator Interface (Decision Rendering)

While the SEP Engine's core lives in a high-performance C++/CUDA environment, its outputs need to be accessible and interpretable to humans and high-level systems. This is achieved via the **external interface layer** (`libsep_api`) and associated UI front-end. The API is exposed as a Crow HTTP server with REST endpoints (as well as a C-callable interface) ⁵. Through these endpoints, one can query the current state of the manifold or request analyses. For example, an endpoint might provide the current predictive gauge or the list of top N patterns by coherence. The UI – which can be a web dashboard or a specialized client application – periodically calls these APIs to fetch fresh data.

Importantly, **the UI does not perform any of the SEP Engine's computations**. It neither computes coherence/stability nor runs simulations – those tasks are exclusively handled by the engine's backend. The UI's role is purely to **render decisions and insights** based on the engine's state. In practical terms, the UI might display: a real-time gauge indicating overall system stability, charts of an identity's coherence over time, or alerts when entropy spikes beyond a threshold. All such content comes directly from reading the Valkey manifold or engine's memory via the API. For instance, if the UI shows that "System Stability = 0.85", this number was calculated by SEP's core (perhaps as an aggregate of pattern stabilities) and stored as a metric, which the UI query simply retrieved. This design follows a classic model-view-controller separation, where SEP's core is the model (and controller for updates), and the UI is the view. The benefit is twofold: **performance isolation** and **conceptual clarity**. Performance-intensive computations (like pattern matching, FFTs from QFH, etc.) are isolated on the GPU machine – the UI is free from heavy processing, which keeps it responsive. Conceptually, the operator can trust that what they see is the authoritative state of the engine, because it's drawn from the engine's own data structures (the Valkey store and memory tiers) rather than any re-computation or approximation on the client side.

Another advantage of this split is safety and auditability. The UI cannot inadvertently alter the engine's state by performing incorrect computations; it only reads state. Decisions (such as trading decisions in a finance context) can thus be made by the human operator or an external logic *with full transparency into the current system state*. The operator interface might highlight, for instance, which pattern identities are currently in the long-term manifold (indicating highly stable patterns) and their coherence values. It could also allow the operator to drill down: because of the manifold's historical memory, the UI could let the user select an identity and scroll back in time through its states (pulling past snapshots from Valkey) to visually inspect how it evolved – essentially a GUI for the manifold's ledger. This kind of introspection is invaluable for explaining the system's outputs, which is often a requirement in fields like finance or medicine (where one might ask "why is the system giving this signal now?"). With SEP's architecture, the answer is available by examining the stored states and metrics trail.

Finally, by splitting **reality tracking vs. decision rendering**, the architecture allows each side to evolve somewhat independently. The engine can be scaled or upgraded (e.g. moving to a larger GPU cluster, or refining the quantum-inspired kernels) without requiring changes to the UI, as long as the API contract remains stable. Conversely, new visualizations or decision support tools can be built on top of the API without altering the core engine. This extensibility is demonstrated by the existence of specialized modules like `sep_blender` (a 3D visualization integration) and `libsep_audio` (an audio-based data ingestion module), which connect to the engine's outputs or inputs without modifying core logic ²⁶. The **UI thus serves as a flexible window onto the SEP Engine's manifold**, enabling users to monitor and interact with the system's emergent patterns in real time, while the engine itself robustly handles the computational heavy lifting behind the scenes.

7. Conclusion

The SEP Engine represents a novel convergence of high-performance computing, information theory, and system design principles. By structuring the problem of pattern discovery as one of maintaining a **manifold of identity states** over time, SEP provides a rigorous and transparent way to track emergent phenomena. The **Valkey database manifold** ensures that every pattern's evolution is time-anchored and queryable, bringing database-like reliability to what is often an ephemeral analytics process. The triad of **coherence, stability, entropy** serves as a set of orthogonal coordinates for measuring pattern "quality," and these in turn define a richly structured, non-Euclidean space in which optimization and inference occur. The engine's **quantum-inspired computational core** (QFH, QBSA) efficiently navigates this space, leveraging CUDA GPU acceleration to update the entire manifold state in real time ²⁷. Thanks to a strictly **deterministic update rule**, the system's behavior is reproducible and theoretically analyzable – a rare trait in complex adaptive systems. This determinism lets us treat the engine as a scientific instrument for studying pattern dynamics: one can rewind, replay, and examine cause and effect with precision ².

For researchers in **quantum computing**, SEP offers an interesting case of quantum concepts (coherence, phase alignment) translated into a classical algorithmic context ³. The notion of a coherence-driven manifold and incremental "register" alignment has parallels to quantum state evolution and measurement ¹⁸, suggesting potential cross-fertilization of ideas (e.g. viewing pattern stability as akin to decoherence time in quantum systems). For experts in **statistical inference**, SEP's metrics resonate with statistical measures (entropy as uncertainty, variance-based stability, etc.), but applied in an online, evolving setting. The engine essentially performs continuous Bayesian-style updating of pattern beliefs, except formulated as deterministic state propagation – an interesting hybrid of statistical filtering and dynamical systems. Those in **non-Euclidean optimization** will recognize that SEP is effectively optimizing pattern representations on a curved information manifold, dealing with trade-offs that are not easily linearized. The success of SEP in finding tradable signals (even yielding measurable *alpha* in live-market tests) stems from this ability to marry rigorous metrics with raw computing power ¹⁶ ¹⁷, traversing an otherwise intractable search space of patterns.

In summary, the SEP Engine's current implementation demonstrates a powerful architecture for reality modeling: **a GPU/Valkey-based reality tracker coupled to a passive but informative UI**, yielding a system that is both **highly operational (for real-time decision support)** and **deeply theoretical (for analysis of emergent phenomena)**. By splitting the "world-state" computation from the "world-view" presentation, SEP adheres to a design that could be described as **measurement-driven computing** – much as in quantum mechanics, where the evolution of a system is separate from the observation of that system. This separation, along with the comprehensive time-anchored state memory, makes the SEP Engine not just

a tool for traders or engineers, but a platform for scientific exploration of pattern emergence and coherence in complex systems. Future work will continue to formalize the mathematics of the SEP manifold and explore its applications, while the robust foundation described here ensures that theoretical insights can be immediately tested and visualized in practice.

References (Code & Documentation)

- SEP Dynamics, *Executive Summary and Technical Overview*, lines 4-11 ¹, 18-24 ³, 54-60 ¹².
- SEP Dynamics Repository, *Predictive Gauge Definition*, lines 56-64 ¹⁷.
- SEP Dynamics Repository, *Proofs of Concept (Reproducibility, etc.)*, lines 62-68 ²².
- **EXPLAIN System Briefing** (Internal Docs), Architecture and Memory, lines 184-192 ¹³, 208-216 ²⁸.
- **Alpha Whitepaper (Draft)**, SEP Metrics Overview, lines 111-118 ²⁹, 119-124 ²⁰.
- **EXPLAIN Technical Briefing** (Internal Docs), External Interface & Visualization, lines 195-203 ⁵.
- SEP Dynamics, *AI-Synthesized Engine and Performance*, lines 61-69 ⁴, 125-132 ²⁷.
- SEP Dynamics, *Mechanic's Analogy (GPU Supercharger)*, lines 61-68 ¹¹.

¹ ² ³ ¹² ¹⁵ ¹⁷ ²² **1_README.md**

https://github.com/SepDynamics/Sep-Dynamics/blob/dd8501d10d68d5b37aafd333812e459d9fff2cd7/backdoor/sepdynamics/1_README.md

⁴ ²⁷ **SEP**

<https://docs.google.com/document/d/14iaUdTGM1ioGepLBf404ru4dYUTF9oYwnEqEs7Ipr3k>

⁵ ⁶ ⁷ ⁸ ⁹ ¹⁰ ¹¹ ¹³ ¹⁴ ²¹ ²³ ²⁴ ²⁵ ²⁶ ²⁸ **EXPLAIN_3_LEVELS.md**

<https://drive.google.com/file/d/1670BUnBA4hETWD6GvIjpj2mBysv-1UCEV>

¹⁶ ¹⁸ ¹⁹ ²⁰ ²⁹ **Alpha_WP.md**

<https://drive.google.com/file/d/1fldBMwvc41gOQZOcdFFf4iZZz5npQQcH>